

NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis

Ben Mildenhall Pratul P. Srinivasan Matthew Tancik
Jonathan T. Barron Ravi Ramamoorthi Ren Ng

UC Berkeley
Google Research
UC San Diego

- Static scene represented as a continuous 5D function:

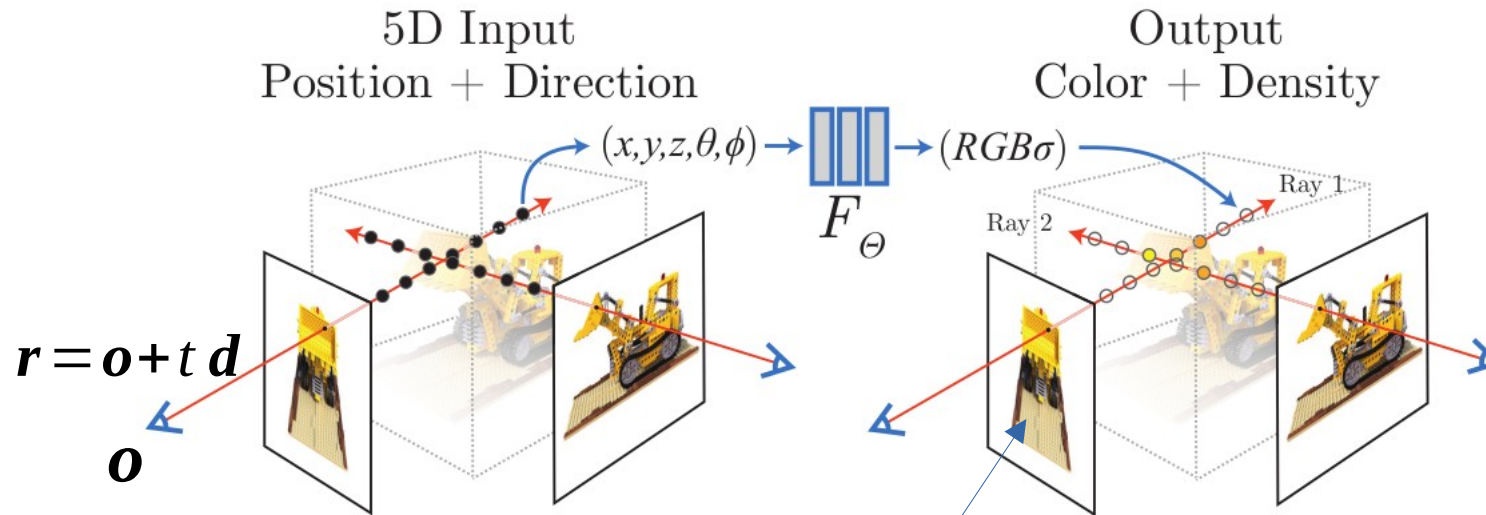
$$F_{\Theta} : ((x, y, z) \stackrel{\text{def}}{=} \mathbf{x}, (\theta, \phi) \stackrel{\text{def}}{=} \mathbf{d}) \rightarrow (\mathbf{c}(\mathbf{x}, \theta, \phi), \sigma(\mathbf{x}))$$

i.e., a *neural radiance field* (NeRF),

- From which an image from any viewpoint can be generated, by
 - marching camera rays ($\mathbf{r} = \mathbf{o} + t \mathbf{d}$) through the image plane to the scene to sample a set of 3D points,
 - Extracting (\mathbf{c}, σ) , then using classical volume rendering techniques to accumulate into a 2D image:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) \mathbf{c}(\mathbf{r}(t), \mathbf{d}) dt, \text{ where } T(t) = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s)) ds\right)$$

Integrate color \times density \times transmittance from near field to far field;
transmittance goes to 0 after some region (i.e., a blocking object) where density \uparrow , meaning
integrated color \times density \times transmittance \rightarrow after that region of blockage.



$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) \mathbf{c}(\mathbf{r}(t), \mathbf{d}) dt$$

- Ray originating from camera generated for each pixel for full 2D render

- Implicit repr. Of continuous 3D shapes via:
 - Signed Distance Functions (SPF) [15, 32]:
output signed (+: outside, -: inside) distance given a 3D point, or
 - Occupancy Fields [11, 27].
- Despite the potential to represent complicated & highres geometry, so far been limited to simple shapes with low geometric complexity,
resulting in oversmoothed renderings

- Given a dense sampling of views, novel views can be reconstructed by simple light field *sampling interpolation techniques* [21, 5, 7]
- With sparser view samplings, novel views are synthesized by predicting **traditional geometry** and **appearance representations** from observed images:
 - Gradient-based mesh optimization based on image representation often difficult due to local minima or poor conditioning of loss
 - Requires template mesh as initialization before optimization (typically unavailable for real-world scenes)

- Volumetric Representations: Able to realistically represent complex shapes and materials & well-suited for gradient-based optimizations & produces less visual artifacts than mesh-based methods.
 - Color **voxel** grids [19, 40, 45]
 - Sampled volumetric representation from a set of input images (point cloud?) → alpha compositing (combining multiple images?) or learned compositing along rays

Highres imagery quadruples time and space complexity because of discrete sampling of 3D space.

- The model: $F_{\Theta}:(\mathbf{x}, \mathbf{d}) \rightarrow (\mathbf{c}, \sigma)$
- We want to generate rays for each pixel on the image plane to compute its color using:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) \mathbf{c}(\mathbf{r}(t), \mathbf{d}) dt, \text{ where } T(t) = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s)) ds\right)$$

which can be numerically estimated using quadrature:


$$\hat{C}(\mathbf{r}) = \sum_{i=1}^N T_i (1 - \exp(-\sigma_i \delta_i)) \mathbf{c}_i \text{ where } T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right), \delta_i = t_{i+1} - t_i$$

→ Trivially Differentiable

Model Formulation: Differentiable Rendering

- Stratified Sampling: Uniformly partition $[t_n, t_f]$ into N evenly spaced bins and draw one sample uniformly at random from each bin, i.e.,:

$$t_i \sim \mathcal{U}\left[t_n + \frac{i-1}{N}(t_f - t_n), t_n + \frac{i}{N}(t_f - t_n)\right]$$




$$\hat{C}(\mathbf{r}) = \sum_{i=1}^N T_i (1 - \exp(-\sigma_i \delta_i)) c_i \text{ where } T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right), \delta_i = t_{i+1} - t_i$$

Model Summary & Training

- Loss: Total squared error btw. The rendered and true pixel colors:

$$\mathcal{L} = \sum_{\mathbf{r} \in \mathcal{R}} \left[\left\| \hat{C}_c(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2 + \left\| \hat{C}_f(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2 \right]$$

- Batch of 4096 rays, each sampled at $N_c = 64$, $N_f = 128$,
- Adam optimizer, learning rate: $5 \times 10^{-4} \rightarrow 5 \times 10^{-5}$ exponential decay
- Optimization for a single scene takes around 100-300k iterations to converge on a single NVIDIA V100 GPU (1-2 days).




- Despite the fact that NN are universal function approximators [14], directly inputting (\mathbf{x}, \mathbf{d}) results in poor estimation of high-frequency variation in color and geometry.
- Rahaman et al. [35] shows deep networks are biased towards learning lower frequency functions, and **mapping inputs to a higher dimensional space using high frequency functions** before passing to network results in better fitting of data that contains high frequency variation.
- So we map (\mathbf{x}, \mathbf{d}) to $(\gamma_{L=10}(\mathbf{x}), \gamma_{L=4}(\mathbf{d}))$ (\mathbf{x} is normalized to lie in $[-1, 1]$, and γ is applied separately to each coordinate value) where
$$\gamma(p) = (\sin(2^0 \pi p), \cos(2^0 \pi p), \dots, \sin(2^{L-1} \pi p), \cos(2^{L-1} \pi p))$$

$$e^{j2^{L-1}\pi p} \Big|_{L=4,10} \xrightarrow{\mathcal{F}} 2\pi\delta(\omega-687), 2\pi\delta(\omega-3\times 10^9)$$

- Densely sampling N points in a ray is inefficient if part of the path is occluded.
- So we optimize two networks: *coarse*— $\hat{C}_c(\mathbf{r})$ — and *fine*— $\hat{C}_f(\mathbf{r})$ —, where the fine network makes sampling decisions based on the output of the coarse network.

Results



Ship

Lego

Microphone

Materials

Ground Truth

NeRF (ours)

LLFF [28]

SRN [42]

NV [24]

| Method | Diffuse Synthetic 360° [41] | | | Realistic Synthetic 360° | | | Real Forward-Facing [28] | | |
|-----------|-----------------------------|--------------|--------------|--------------------------|--------------|--------------|--------------------------|--------------|--------------|
| | PSNR↑ | SSIM↑ | LPIPS↓ | PSNR↑ | SSIM↑ | LPIPS↓ | PSNR↑ | SSIM↑ | LPIPS↓ |
| SRN [42] | 33.20 | 0.963 | 0.073 | 22.26 | 0.846 | 0.170 | 22.84 | 0.668 | 0.378 |
| NV [24] | 29.62 | 0.929 | 0.099 | 26.05 | 0.893 | 0.160 | - | - | - |
| LLFF [28] | 34.38 | 0.985 | 0.048 | 24.88 | 0.911 | 0.114 | 24.13 | 0.798 | 0.212 |
| Ours | 40.15 | 0.991 | 0.023 | 31.01 | 0.947 | 0.081 | 26.50 | 0.811 | 0.250 |

- Diffuse Synthetic: Lambertian materials only
- Metrics:
 - PSNR (Peak Signal-to-Noise Ratio): $10 \log MAX^2 / MSE$
 - SSIM (Structural Similarity Index Measure): evaluation based on brightness, contrast, and structure; $SSIM \in [0, 1]$
 - LPIPS [50]: perceptual loss; DL-based

- A full training of images annotated with camera parameters is required to produce images from novel angles.
- Poor performance when camera strays from center of object?
- Poor mesh restoration?